

# Making Decisions with Reinforcement Learning

Todd Sierens<sup>1,2</sup>

<sup>1</sup>Perimeter Institute for Theoretical Physics

<sup>2</sup>University of Waterloo

May 5, 2016

# Outline


- 1 Machine Learning
- 2 Neural Networks
  - Simple Perceptron
  - Multilayer Perceptron
  - Convolutional Neural Network
- 3 Reinforcement Learning
  - Temporal Difference Method
  - Stochastic Gradient Descent
- 4 Results
- 5 Future Directions

# Supervised Learning vs Reinforcement Learning

- Supervised Learning
  - Have access to a large set of data with known desired results.
  - Adjust model parameters to minimize an objective function.
- Reinforcement Learning
  - Have access to an environment that can be modeled.
  - Typically a reward function is used as a signal for how to adjust model parameters.
  - Board games present a natural environment that is easily modeled and can provide a reward whenever a game completes.

# Simple Perceptron

Input



Hidden




Output







# Simple Perceptron



# Simple Perceptron




# Multilayer Perceptron





# Convolutional Neural Network






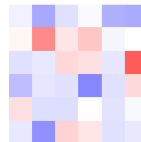
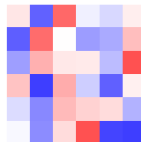
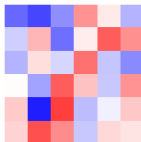
# Value Network

- A value network is a neural network that evaluates an environment and determines a value to associate to it.
- In the case of board games, a value network can be used to determine the probability of winning from any given position.
- Here is an example of a multi-layered neural network in action.
- The network takes a Tic Tac Toe board as input, and through a series of node activations the network outputs a prediction on who will win, in this case: the red player.

Tic Tac Toe Position



Network Node Activations



Predicted Winner



# Reinforcement Learning

## ■ TD( $\lambda$ ) Equation

- When a terminal state occurs, the reward propagates to previous states and the targets are updated

$$T_n = (1 - \lambda) V_n + \lambda r$$


# Reinforcement Learning


## ■ TD( $\lambda$ ) Equation


- When a terminal state occurs, the reward propagates to previous states and the targets are updated


$$T_n = \sum_{n=0}^{N_0} (1 - \lambda) \lambda^n V_n + \lambda^{N_0} r$$


# Stochastic Gradient Descent













# Future Directions

- Apply to more games
  - chess, hex, backgammon, go, video games, etc.

Route finding app

- Traffic prediction analysis
- Control (robots, self-driving cars, etc.)
- Prediction (stock prices, sports betting, etc.)
- Acquire large data sets and compare with supervised learning